

END-TO-END DEVOPS AUTOMATION PROJECT

A Production-Ready CI/CD Pipeline with Spring Boot, Docker,
Kubernetes, Terraform, Jenkins, Prometheus, Grafana, and AWS EKS

FINAL REPORT

Submitted in partial fulfillment of the requirements for the degree
of Bachelor of Technology in Computer Science & Engineering

Submitted by:

Kakumani Phaneendra

Roll Number: 12418953

Under the guidance of:

Ajay Kumar Badhan

Department of Computer Science & Engineering

Lovely Professional University

Phagwara, Punjab

INT377: Cloud Computing and DevOps Essentials

May 19, 2026

Certificate

This is to certify that the project titled “**End-to-End DevOps Automation Project**” is a bonafide work carried out by **Kakumani Phaneendra** bearing roll number **12418953** in partial fulfillment of the requirements for the degree of **Bachelor of Technology in Computer Science & Engineering** at **Lovely Professional University**. The project demonstrates a complete, production-grade DevOps pipeline incorporating continuous integration, continuous deployment, containerization, orchestration, infrastructure as code, and observability.

The work is original and has not been submitted for any other degree or diploma.

Date: May 19, 2026

Place: Punjab

Declaration

I, **Kakumani Phaneendra**, hereby declare that this project report entitled “**End-to-End DevOps Automation Project**” is my own original work. All sources of information, code snippets, configuration files, and third-party references have been duly acknowledged. This work has not been submitted previously for any other degree or diploma.

Date: May 19, 2026

Place: Phagwara, Punjab

Signature: Kakumani Phaneendra

Acknowledgement

I express my sincere gratitude to my project guide, **Ajay Kumar Badhan**, for their invaluable guidance, technical insights, and continuous encouragement throughout this project. Their deep understanding of DevOps methodologies and cloud-native technologies greatly enriched this work.

I extend my heartfelt thanks to the Head of the Department, **Pushpendra Kumar Pateriya**, for providing the necessary infrastructure and support. I also acknowledge the contributions of the faculty and staff of the Computer Science & Engineering department for their cooperation and motivation.

Special thanks to the open-source communities behind Spring Boot, Jenkins, Docker, Kubernetes, Terraform, Prometheus, and Grafana for making enterprise-grade DevOps accessible to everyone.

Finally, I thank my family and friends for their unwavering support and patience.

Kakumani Phaneendra

Abstract

This project presents a complete end-to-end DevOps automation solution that integrates modern development, deployment, and monitoring practices. A Spring Boot REST API application with an interactive static frontend is containerized using Docker (multi-stage builds) and orchestrated on Amazon EKS (Kubernetes). Terraform is used as Infrastructure as Code (IaC) to provision all AWS resources – VPC, EKS cluster, ECR repository, IAM roles – in a repeatable and version-controlled manner. A Jenkins CI/CD pipeline automates the entire workflow: source code checkout from GitHub, Maven build and unit testing, Docker image creation, image push to Amazon ECR, Terraform

infrastructure validation/apply, Kubernetes deployment with rolling updates, and smoke testing.

Monitoring is implemented using Prometheus (metrics collection from Spring Boot Actuator) and Grafana (real-time dashboards). The project includes a rich frontend dashboard with a CI/CD pipeline simulator, live console logs, interactive cluster map, and gamified debugging games (Snake, Bug Squasher) to demonstrate DevOps concepts engagingly.

The system eliminates manual deployment errors, reduces time-to-market from hours to minutes, ensures environment consistency, provides automatic scaling and self-healing, and offers full observability. All code, configurations, and documentation are version-controlled in GitHub. This report provides a complete, reproducible, step-by-step guide suitable for production environments, academic learning, and professional portfolio demonstration.

Keywords: DevOps, CI/CD, Jenkins, Docker, Kubernetes, Terraform, AWS EKS, Spring Boot, Prometheus, Grafana, Infrastructure as Code, Automation, Rolling Updates, Multi-stage Builds

Table of Contents

1. [Introduction](#)
2. [Problem Statement](#)
3. [Objectives](#)
4. [Technology Stack](#)
5. [System Requirements](#)
6. [Project Architecture](#)
7. [Step-by-Step Implementation](#)
 - 7.1 [Initial Setup & Version Control](#)
 - 7.2 [Backend Development \(Spring Boot\)](#)
 - 7.3 [Frontend Dashboard & Gamification](#)
 - 7.4 [Dockerization \(Multi-Stage Build\)](#)
 - 7.5 [Kubernetes Manifests](#)
 - 7.6 [Terraform Infrastructure \(AWS\)](#)
 - 7.7 [Jenkins CI/CD Pipeline](#)
 - 7.8 [Monitoring with Prometheus & Grafana](#)
8. [Source Code Explanation](#)
9. [Deployment Process](#)
10. [CI/CD Pipeline Explanation](#)

11. [Infrastructure Setup](#)
 12. [Testing Strategy & Results](#)
 13. [Debugging & Troubleshooting Guide](#)
 14. [Security Best Practices](#)
 15. [Performance Optimization Techniques](#)
 16. [Challenges Faced & Solutions](#)
 17. [Future Enhancements](#)
 18. [Conclusion](#)
- [References](#)
- [Appendix A: Full Command Reference](#)
 - [Appendix B: Screenshot Placeholders](#)
 - [Appendix C: Environment Variables & Secrets](#)

1. Introduction

DevOps is a cultural and technical movement that bridges the gap between software development (Dev) and IT operations (Ops). It emphasizes automation, continuous integration (CI), continuous delivery/deployment (CD), infrastructure as code (IaC), monitoring, and collaboration. Traditional software deployment often involves manual, error-prone steps: developers build artifacts locally, operations teams manually configure servers, and inconsistencies lead to “it works on my machine” problems. This results in slow release cycles, environment drift, difficulty scaling, and limited observability.

This project implements a fully automated, production-grade DevOps pipeline for a cloud-native web application consisting of a Spring Boot backend and a static frontend with interactive components. The pipeline automatically:

- Builds and tests the application on every code push.
- Creates a minimal Docker image using multi-stage builds.
- Stores the image in Amazon ECR (private registry).
- Provisions or updates AWS infrastructure (VPC, EKS cluster, node groups) using Terraform.
- Deploys the new image to a Kubernetes cluster (EKS) with rolling updates.
- Monitors application health and performance using Prometheus and Grafana.

The project demonstrates key DevOps principles in action:

Principle	Implementation
-----------	----------------

Continuous Integration	Jenkins triggers on GitHub push; Maven compiles and runs JUnit tests.
Continuous Deployment	After successful CI, the pipeline automatically deploys to Kubernetes.
Infrastructure as Code	All AWS resources defined in Terraform scripts.
Containerization	Docker packaging ensures environment consistency from dev to prod.
Orchestration	Kubernetes manages scaling, self-healing, and rolling updates.
Observability	Prometheus scrapes metrics; Grafana visualizes them.

The report is structured as a complete implementation guide – from setting up the development environment to deploying and monitoring the application in the cloud. Every command, configuration file, and decision is explained in detail, making it suitable for both academic submission and real-world reference.

Reference documentation: [Spring Boot Reference Guide](#) | [Docker Docs](#) | [Kubernetes Docs](#) | [Terraform Docs](#)

2. Problem Statement

In many organizations, software deployment remains a manual, fragile, and time-consuming process. Specific pain points include:

- **Environment Inconsistency** – Code works on a developer’s machine but fails in production due to differing OS versions, library versions, or configuration settings.
- **Slow Release Cycles** – Manual builds, tests, and deployments can take hours or even days, delaying feature delivery and bug fixes.
- **Scaling Challenges** – Without automation, handling traffic spikes requires manual provisioning of new servers, leading to downtime or degraded performance.

- **Infrastructure Drift** – Over time, servers become unique snowflakes because of ad-hoc changes (e.g., configuration tweaks, package installations). This makes it impossible to recreate the environment reliably.
- **Limited Observability** – Teams often lack real-time insight into application health, performance metrics, and error rates, making proactive issue detection difficult.
- **No Rollback Strategy** – When a deployment fails, reverting to a previous stable version is often complex and slow.

These problems increase operational costs, reduce developer productivity, cause frequent production incidents, and hinder business agility. This project solves them by automating the entire lifecycle and providing a unified observability platform.

3. Objectives

Primary Objectives

#	Objective	How It Is Achieved
1	Automate Build & Test	Jenkins pipeline runs <code>mvn clean package</code> and JUnit tests on every push.
2	Containerize the Application	Multi-stage Dockerfile produces a portable image.
3	Provision AWS Infrastructure as Code	Terraform scripts define VPC, EKS cluster, ECR, IAM roles.
4	Orchestrate with Kubernetes	Deploy to EKS; use Deployments, Services, rolling updates.
5	Implement CI/CD Pipeline	End-to-end automation from code commit to production deployment.
6	Set Up Monitoring & Observability	Prometheus collects metrics; Grafana provides dashboards.
7	Provide Complete Documentation	Detailed step-by-step guide with commands, configs, and explanations.

Secondary Objectives

- Multi-stage Docker builds to minimize image size (security & performance).
- Liveness & Readiness probes in Kubernetes for self-healing.
- AWS Load Balancer for external access.
- Interactive frontend with gamified elements to demonstrate DevOps concepts.
- Security best practices (non-root containers, least privilege IAM, secrets).

4. Technology Stack

Category	Technology	Version	Purpose	Justification
Backend Framework	Spring Boot	3.2.5	REST API development , dependency injection	Latest stable with Java 21 support; excellent actuator/metrics integration.
Programming Language	Java (OpenJDK)	21 LTS	Application runtime	Long-term support, modern features, performance improvements.
Build Tool	Maven	3.9+	Dependency management , compilation, testing	Standard for Spring Boot; declarative pom.xml.
Frontend	HTML5, CSS3, JavaScript	-	Dashboard, games, real-time status	Works without additional build steps; served as static content.
Containerization	Docker	24.0+	Image creation, packaging	Industry standard; multi-stage builds supported.

Container Registry	Amazon ECR	-	Private Docker image storage	Integrated with AWS IAM; scanning on push.
Orchestration	Kubernetes (EKS)	1.28	Pod management , scaling, rolling updates	Managed service reduces operational overhead; cloud-native.
Local Orchestration	Minikube	latest	Local testing environment	Lightweight, simulates Kubernetes locally.
Infrastructure as Code	Terraform	1.6+	AWS resource provisioning	Declarative, version-controlled , supports multiple providers.
CI/CD Server	Jenkins	2.426+	Pipeline automation	Highly extensible, declarative pipeline support.
Version Control	Git / GitHub	-	Source code management	Webhooks enable automatic pipeline triggers.
Cloud Provider	AWS	-	Compute, networking, registry	Market leader; EKS, ECR, VPC, IAM.
Monitoring	Prometheus	2.48+	Metrics collection	Pull-based, integrates with Spring Boot Actuator.
Visualization	Grafana	10.2+	Dashboards	Powerful, customizable,

				connects to Prometheus.
Testing	JUnit 5, TestRestTemplate	5.10	Unit & integration tests	Spring Boot starter test.

Version Rationale:

- **Java 21** – LTS, virtual threads, pattern matching.
- **Spring Boot 3.2.5** – Compatible with Java 21, improved observability (Micrometer).
- **Kubernetes 1.28** – Latest stable, enhanced security features.
- **Terraform 5.x AWS provider** – Full EKS support, stable.
- **Jenkins 2.426** – Declarative pipeline mature, plugin ecosystem.

5. System Requirements

5.1 Development Environment (Local Machine)

Component	Minimum	Recommended
CPU	4 cores	8 cores (for Minikube + Jenkins + IDE)
RAM	8 GB	16 GB
Disk Space	50 GB SSD	100 GB SSD
Operating System	Ubuntu 22.04 / macOS 13 / Windows 10 Pro (with WSL2)	Ubuntu 22.04 LTS
Network	10 Mbps broadband	50+ Mbps

5.2 Required Software (Development Machine)

All tools must be installed before starting the project. Below are installation commands for Ubuntu (similar for other OSes).

Tool	Version	Installation Command (Ubuntu)
Git	2.40+	<code>sudo apt install git -y</code>
Docker	24.0+	<code>` curl -fsSL https://get.docker.com</code>
kubect1	1.28+	<code>curl -LO "https://dl.k8s.io/release/v1.28.0/bin/linux/amd64/kubect1" && chmod +x kubect1 && sudo mv kubect1 /usr/local/bin/</code>
Minikube	latest	<code>curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64 && sudo install minikube-linux-amd64 /usr/local/bin/minikube</code>
Terraform	1.6+	Instructions from HashiCorp
AWS CLI	2.13+	<code>curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip" && unzip awscliv2.zip && sudo ./aws/install</code>
Java 21 JDK	21	<code>sudo apt install openjdk-21-jdk -y</code>
Maven	3.9+	<code>sudo apt install maven -y</code>
Jenkins (container)	2.426+	<code>docker run -d --name jenkins -p 8080:8080 -p 50000:50000 -v jenkins_home:/var/jenkins_home -v /var/run/docker.sock:/var/run/docker.sock jenkins/jenkins:lts</code>

5.3 AWS Production Infrastructure (EKS)

Resource	Specification	Purpose
EKS Cluster	1 control plane (managed)	Kubernetes control plane
Worker Nodes	2 × t3.medium (min), up to 3	Run application pods


VPC	Custom, 2 public subnets (us-east-1a, us-east-1b)	Network isolation, Load Balancer placement
Internet Gateway	1	Public internet access for nodes and LB
ECR Repository	1 private	Store Docker images
IAM Roles	2 (cluster, node group) + Jenkins user	Permissions for EKS and CI/CD

Estimated Monthly Cost (us-east-1):

- EKS cluster: $\$0.10/\text{hr} \times 730 = \73
- t3.medium EC2 (2 nodes): $\$0.0416 \times 2 \times 730 \approx \60.7
- NLB: $\$0.025 \times 730 \approx \18.25
- **Total $\approx$ \$152/month** (reduce by destroying after demo or using spot instances).

6. Project Architecture

6.1 High-Level Architecture Diagram



6.2 Component Descriptions

Component	Role	Interaction
GitHub	Source code repository	Stores application code, K8s manifests, Terraform scripts, Jenkinsfile. Webhook triggers Jenkins.
Jenkins	CI/CD automation server	Executes pipeline stages, builds JAR, creates Docker image, pushes to ECR, runs Terraform, deploys to K8s.
Docker	Containerization	Packages Spring Boot app and static files into a portable image.
Amazon ECR	Private Docker registry	Stores images securely; Jenkins pushes each build.
Terraform	Infrastructure as Code	Provisions AWS VPC, subnets, IGW, EKS cluster, node group, ECR, IAM roles.
AWS EKS	Managed Kubernetes	Runs the application pods; manages scaling, self-healing, rolling updates.
Prometheus	Metrics collection	Scrapes /actuator/prometheus from app pods; also collects node metrics.
Grafana	Visualization	Displays dashboards (CPU, memory, request rates) using Prometheus data.
AWS Load Balancer	Traffic ingress	Exposes the application to the internet; routes to healthy pods.

6.3 Data Flow (End-to-End)

1. **Code Push** – Developer runs `git push origin main`.
2. **Webhook** – GitHub sends a POST request to Jenkins' webhook URL.
3. **Pipeline Trigger** – Jenkins starts a new build.
4. **Checkout** – Jenkins clones the repository.
5. **Build & Test** – `mvn clean package` runs unit tests; if any test fails, pipeline stops.
6. **Docker Build** – Multi-stage Dockerfile creates an image tagged with build number.

7. **Push to ECR** – Jenkins logs into AWS ECR and pushes the image.
8. **Terraform** – Terraform plans and applies any infrastructure changes (idempotent).
9. **Kubernetes Deploy** – `kubectl set image` updates the deployment to use the new image; Kubernetes performs a rolling update.
10. **Smoke Test** – Jenkins curls the load balancer endpoint to verify the app is healthy.
11. **Monitoring** – Prometheus continuously scrapes metrics; Grafana displays them.

7. Step-by-Step Implementation

7.1 Initial Setup & Version Control

7.1.1 Install Prerequisites (Ubuntu 22.04)

```
# Update system
sudo apt update && sudo apt upgrade -y

# Install Git
sudo apt install git -y
git --version
# Output: git version 2.40.1

# Install Java 21 JDK
sudo apt install openjdk-21-jdk -y
java -version
# Output: openjdk version "21.0.2" 2024-01-16 LTS

# Install Maven
sudo apt install maven -y
mvn -version
# Output: Apache Maven 3.9.4

# Install Docker
curl -fsSL https://get.docker.com -o get-docker.sh
sudo sh get-docker.sh
sudo usermod -aG docker $USER
newgrp docker
docker --version
# Output: Docker version 24.0.7
```

```
# Install kubectl
curl -LO "https://dl.k8s.io/release/v1.28.0/bin/linux/amd64/kubectl"
chmod +x kubectl
sudo mv kubectl /usr/local/bin/
kubectl version --client
# Output: Client Version: v1.28.0

# Install Terraform
wget -O- https://apt.releases.hashicorp.com/gpg | gpg --dearmor | sudo tee
/usr/share/keyrings/hashicorp-archive-keyring.gpg
echo "deb [signed-by=/usr/share/keyrings/hashicorp-archive-
keyring.gpg] https://apt.releases.hashicorp.com $(lsb_release -cs) main" |
sudo tee /etc/apt/sources.list.d/hashicorp.list
sudo apt update && sudo apt install terraform -y
terraform version
# Output: Terraform v1.6.5

# Install AWS CLI
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o
"awscliv2.zip"
unzip awscliv2.zip
sudo ./aws/install
aws --version
# Output: aws-cli/2.13.0
```

Reference links for installation:

- [Git downloads](#)
- [Docker Engine installation](#)
- [kubectl installation](#)
- [Terraform installation](#)
- [AWS CLI installation](#)
- [Java 21 JDK](#)
- [Maven installation](#)

7.1.2 Configure AWS CLI

```
aws configure
# AWS Access Key ID: [Enter your IAM user access key]
# AWS Secret Access Key: [Enter your secret key]
```

```
# Default region: us-east-1
# Default output format: json

# Verify identity
aws sts get-caller-identity
# Output: {"UserId": "AIDA...", "Account": "123456789012", "Arn":
"arn:aws:iam::123456789012:user/yourname"}
```

After configuration, you can verify by visiting the [AWS Management Console](#).

7.1.3 Create GitHub Repository and Push Initial Code

Create a new repository at [GitHub](#).

Use the URL: <https://github.com/yourusername/devops-automation-project>

```
# Configure Git user
git config --global user.name "Your Name"
git config --global user.email "your.email@example.com"

# Clone the repository
git clone https://github.com/yourusername/devops-automation-project.git
cd devops-automation-project

# Generate Spring Boot project using Spring Initializr (via curl)
curl https://start.spring.io/starter.zip \
  -d dependencies=web,actuator,prometheus \
  -d javaVersion=21 \
  -d groupId=com.example \
  -d artifactId=devops-app \
  -o devops-app.zip
unzip devops-app.zip
rm devops-app.zip

# Add all files to Git
git add .
git commit -m "Initial Spring Boot project with Actuator and
Prometheus"
git push -u origin main
```


GitHub repository URL: <https://github.com/yourusername/devops-automation-project>

Spring Initializr URL: <https://start.spring.io/>

7.2 Backend Development (Spring Boot)

7.2.1 Project Structure After Development

```
devops-app/
├── pom.xml
├── src/
│   ├── main/
│   │   ├── java/com/example/devopsapp/
│   │   │   ├── DevopsAppApplication.java
│   │   │   ├── controller/
│   │   │   │   └── AppController.java
│   │   │   ├── service/
│   │   │   │   └── AppService.java
│   │   │   ├── model/
│   │   │   │   ├── AppMetadata.java
│   │   │   │   └── GitMetadata.java
│   │   └── resources/
│   │       ├── application.properties
│   │       ├── static/
│   │       │   ├── index.html
│   │       │   ├── status.html
│   │       │   ├── css/style.css
│   │       │   └── js/
│   │       │       ├── snake.js
│   │       │       └── bug-squasher.js
│   └── test/java/com/example/devopsapp/
│       └── DevopsAppApplicationTests.java
├── Dockerfile (to be added)
├── Jenkinsfile (to be added)
├── kubernetes/ (to be added)
├── terraform/ (to be added)
└── monitoring/ (to be added)
```

7.2.2 Key Backend Files (Detailed Explanation)

pom.xml – Dependencies and Build Configuration

```
<?xml version="1.0" encoding="UTF-8"?>
<project xmlns="http://maven.apache.org/POM/4.0.0" ...>
  <modelVersion>4.0.0</modelVersion>
  <parent>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-parent</artifactId>
    <version>3.2.5</version>
  </parent>
  <groupId>com.example</groupId>
  <artifactId>devops-app</artifactId>
  <version>0.0.1-SNAPSHOT</version>
  <properties>
    <java.version>21</java.version>
  </properties>
  <dependencies>
    <!-- Spring Web MVC -->
    <dependency>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter-web</artifactId>
    </dependency>
    <!-- Actuator for health, metrics, info endpoints -->
    <dependency>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter-actuator</artifactId>
    </dependency>
    <!-- Prometheus registry for metrics export -->
    <dependency>
      <groupId>io.micrometer</groupId>
      <artifactId>micrometer-registry-prometheus</artifactId>
    </dependency>
    <!-- Testing -->
    <dependency>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter-test</artifactId>
      <scope>test</scope>
    </dependency>
  </dependencies>
  <build>
    <plugins>
```

```

        <plugin>
            <groupId>org.springframework.boot</groupId>
            <artifactId>spring-boot-maven-plugin</artifactId>
        </plugin>
    </plugins>
</build>
</project>

```

Why these dependencies?

- `spring-boot-starter-web` – REST APIs.
- `spring-boot-starter-actuator` – Production-ready endpoints (/health, /info, /metrics).
- `micrometer-registry-prometheus` – Exposes metrics in Prometheus format at /actuator/prometheus.

application.properties – Configuration

```

server.port=8081
management.endpoints.web.exposure.include=*
management.endpoint.health.show-details=always
management.info.env.enabled=true
management.metrics.export.prometheus.enabled=true
management.metrics.tags.application=devops-app
info.app.name=devops-app
info.app.version=0.0.1-SNAPSHOT
info.app.description=Spring Boot Cloud-Native Application

```

- `server.port=8081` – The app runs on port 8081 (instead of default 8080).
- `management.endpoints.web.exposure.include=*` – Exposes all actuator endpoints (health, info, prometheus, etc.) over HTTP.
- `management.endpoint.health.show-details=always` – Shows detailed health info (disk space, etc.).
- `management.info.env.enabled=true` – Includes environment properties in /info.
- `management.metrics.export.prometheus.enabled=true` – Explicitly enables Prometheus export.
- `management.metrics.tags.application=devops-app` – Adds an application label to every metric.

AppController.java – REST Endpoints with Content Negotiation

```

package com.example.devopsapp.controller;

import com.example.devopsapp.model.AppMetadata;
import com.example.devopsapp.model.GitMetadata;
import com.example.devopsapp.service.AppService;
import org.springframework.http.MediaType;
import org.springframework.http.ResponseEntity;
import org.springframework.web.bind.annotation.*;

import java.util.Map;

@RestController
public class AppController {

    private final AppService appService;

    public AppController(AppService appService) {
        this.appService = appService;
    }

    // Health endpoint - returns JSON or HTML based on Accept header
    @GetMapping(value = "/health", produces =
{MediaType.APPLICATION_JSON_VALUE, MediaType.TEXT_HTML_VALUE})
    public ResponseEntity<?> health(@RequestHeader(value = "Accept",
defaultValue = "") String accept) {
        Map<String, String> health = appService.getHealthStatus();
        if (accept.contains("text/html")) {
            return ResponseEntity.ok(appService.wrapInHtml("System
Health",
                "Status: <b>" + health.get("status") + "</b><br>" +
health.get("message"), "heart-pulse"));
        }
        return ResponseEntity.ok(health);
    }

    // Greeting endpoint
    @GetMapping(value = "/hello", produces =
{MediaType.TEXT_PLAIN_VALUE, MediaType.TEXT_HTML_VALUE})
    public ResponseEntity<String> hello(@RequestParam(required =
false) String name,
                                     @RequestHeader(value =
"Accept", defaultValue = "") String accept) {

```

```

        String greeting = appService.getGreeting(name);
        if (accept.contains("text/html")) {
            return
ResponseEntity.ok(appService.wrapInHtml("Greeting", greeting,
"message-square"));
        }
        return ResponseEntity.ok(greeting);
    }

    // Application metadata
    @GetMapping(value = "/info", produces =
{MediaType.APPLICATION_JSON_VALUE, MediaType.TEXT_HTML_VALUE})
    public ResponseEntity<?> info(@RequestHeader(value = "Accept",
defaultValue = "") String accept) {
        AppMetadata info = appService.getAppMetadata();
        if (accept.contains("text/html")) {
            return
ResponseEntity.ok(appService.wrapInHtml("Application Info",
                "Name: <b>" + info.name() + "</b><br>Version: " +
info.version() + "<br>Java: " + info.javaVersion(), "info"));
        }
        return ResponseEntity.ok(info);
    }

    // Git metadata (branch, commit)
    @GetMapping(value = "/git", produces =
{MediaType.APPLICATION_JSON_VALUE, MediaType.TEXT_HTML_VALUE})
    public ResponseEntity<?> git(@RequestHeader(value = "Accept",
defaultValue = "") String accept) {
        GitMetadata git = appService.getGitInfo();
        if (accept.contains("text/html")) {
            return ResponseEntity.ok(appService.wrapInHtml("Git
Repository",
                "Branch: <b>" + git.branch() + "</b><br>Commit:
<code>" + git.commitId() + "</code><br>Message: <i>" +
git.commitMessage() + "</i>", "git-branch"));
        }
        return ResponseEntity.ok(git);
    }

    // Metrics viewer page
    @GetMapping(value = "/metrics-view", produces =

```

```

MediaType.TEXT_HTML_VALUE)
    public String metricsView() {
        return appService.wrapInHtml("System Metrics",
            "Live Prometheus metrics are being collected.<br><br><a
href='/actuator/prometheus' style='color: #00f2fe;'>View Raw
Data</a>", "bar-chart-3");
    }
}

```

Explanation of Content Negotiation: The produces attribute allows the same endpoint to return either JSON (for API clients) or HTML (for browsers). The Accept header determines the response format. This makes the API developer-friendly and also directly usable in a browser.

AppService.java – Business Logic

```

package com.example.devopsapp.service;

import com.example.devopsapp.model.AppMetadata;
import com.example.devopsapp.model.GitMetadata;
import org.springframework.stereotype.Service;

import java.util.Map;

@Service
public class AppService {

    public Map<String, String> getHealthStatus() {
        return Map.of("status", "UP", "message", "Application is
healthy");
    }

    public String getGreeting(String name) {
        return "Hello, " + (name != null ? name : "DevOps") + "!";
    }

    public AppMetadata getAppMetadata() {
        return new AppMetadata(
            "devops-app",
            "0.0.1-SNAPSHOT",
            "Spring Boot Cloud-Native Application",
            System.getProperty("java.version")
        );
    }
}

```

```

    );
}

public GitMetadata getGitInfo() {
    // Execute git commands to get current branch and commit
    String branch = executeCommand("git rev-parse --abbrev-ref
HEAD");
    String commitId = executeCommand("git rev-parse --short
HEAD");
    String message = executeCommand("git log -1 --pretty=%B");

    return new GitMetadata(
        branch != null ? branch : "unknown",
        commitId != null ? commitId : "unknown",
        message != null ? message.trim() : "unknown"
    );
}

private String executeCommand(String command) {
    try {
        Process process = Runtime.getRuntime().exec(command);
        java.util.Scanner s = new
java.util.Scanner(process.getInputStream()).useDelimiter("\\A");
        return s.hasNext() ? s.next().trim() : null;
    } catch (Exception e) {
        return null;
    }
}

// HTML wrapper for endpoints (creates a styled page)
public String wrapInHtml(String title, String content, String
icon) {
    return ""
        <!DOCTYPE html>
        <html>
        <head><title>%s | DevOps App</title>
        <style>/* CSS styles as in style.css */</style>
        </head>
        <body>... HTML structure ...</body>
        </html>
        """.formatted(title, icon, title, content);
}

```

```
}
```

Why execute git commands? In a development environment, this retrieves the actual Git metadata. In production (inside a container), the `.git` folder is not present, so it returns "unknown". For production, we could use the `git-commit-id-plugin` to generate a `git.properties` file at build time.

Model Records:

```
// AppMetadata.java
package com.example.devopsapp.model;
public record AppMetadata(String name, String version, String
description, String javaVersion) {}

// GitMetadata.java
package com.example.devopsapp.model;
public record GitMetadata(String branch, String commitId, String
commitMessage) {}
```

Records (Java 14+) provide immutable data carriers with built-in constructor, getters, equals, hashCode, and toString.

7.2.3 Build and Test Locally

```
# Build the project
mvn clean package
# Output: [INFO] Building jar: target/devops-app-0.0.1-SNAPSHOT.jar

# Run the application
java -jar target/devops-app-0.0.1-SNAPSHOT.jar
# Output: Started DevopsAppApplication in 2.5 seconds

# Test endpoints (in another terminal)
curl http://localhost:8081/health
# {"status":"UP","message":"Application is healthy"}

curl "http://localhost:8081/hello?name=Jenkins"
# Hello, Jenkins!

curl http://localhost:8081/info
# {"name":"devops-app","version":"0.0.1-
```



```
SNAPSHOT", "description": "Spring Boot Cloud-Native  
Application", "javaVersion": "21.0.2"}
```

```
curl http://localhost:8081/git  
# {"branch": "main", "commitId": "a1b2c3d", "commitMessage": "Initial  
commit"}  
  
curl http://localhost:8081/actuator/health  
#  
{"status": "UP", "components": {"diskSpace": {"status": "UP", "details": {  
..}}}}
```

7.3 Frontend Dashboard & Gamification

The frontend consists of static files placed in `src/main/resources/static/`. Spring Boot automatically serves them from the root context path (`/`). The main pages are:

- **index.html** – DevOps Control Center with pipeline simulator, live console, interactive cluster map, and game selector.
- **status.html** – Real-time system health monitor that polls `/health`, `/git`, `/info` every 5 seconds.
- **css/style.css** – Dark theme, glassmorphism, animations, responsive grid.
- **js/snake.js** – Classic Snake game (keyboard controls) – metaphor for orchestration.
- **js/bug-squasher.js** – Click-to-squash bugs game – metaphor for debugging.

index.html Key Sections:

- **Pipeline Simulator:** When user clicks “Deploy to Production”, JavaScript animates five steps (Build, Unit Test, Security Scan, Push ECR, K8s Deploy) with a 1.5-second delay each, updating the UI and adding log messages to the console.
- **Console Log:** Displays real-time simulated logs (VPC established, cluster ready, health checks). Random logs are added every 8 seconds.
- **Cluster Map:** Four node icons (Node-A, B, C, D) with different status colors. Clicking shows a modal with node details (region, instance type, IP, uptime).
- **Game Zone:** Canvas element with Snake or Bug Squasher game. Users can switch games via dropdown.

snake.js Highlights:

```

const gridSize = 20;
const tileCount = canvas.width / gridSize;
let snake = [{x:10, y:10}, {x:10, y:11}, {x:10, y:12}];
let dx = 0, dy = -1; // moving up initially

function drawGame() {
    clearCanvas();
    moveSnake();
    checkCollision();
    drawFood();
    drawSnake();
}

```

status.html – Real-Time Monitor: Uses `fetch()` to call `/health`, `/git`, `/info` every 5 seconds and updates the DOM. Displays status, branch, commit, app name, version, Java version.

7.4 Dockerization (Multi-Stage Build)

Create a Dockerfile in the project root:

```

# Stage 1: Build (uses full Maven + JDK)
FROM maven:3.9-eclipse-temurin-21 AS builder
WORKDIR /app
COPY pom.xml .
RUN mvn dependency:go-offline    # Download dependencies to cache
them
COPY src ./src
RUN mvn clean package -DskipTests

# Stage 2: Runtime (only JRE Alpine)
FROM eclipse-temurin:21-jre-alpine
WORKDIR /app
# Install curl for healthcheck (production readiness)
RUN apk add --no-cache curl
# Create a non-root user for security
RUN addgroup -S appgroup && adduser -S appuser -G appgroup
COPY --from=builder /app/target/*.jar app.jar
USER appuser
EXPOSE 8081
ENTRYPOINT ["java", "-jar", "app.jar"]

```

Why multi-stage?

- Stage 1 includes Maven and the full JDK (large ~600 MB).
- Stage 2 only includes the JRE (smaller ~180 MB).
- Final image does not contain build tools, reducing attack surface and storage/transfer costs.
- The `apk add --no-cache curl` line ensures the Docker Compose healthcheck can work correctly.

Build and test the Docker image:

```
docker build -t devops-app:latest .
docker images | grep devops-app
# devops-app    latest    abc123    2 minutes ago    182MB

docker run -d -p 8081:8081 --name devops-test devops-app:latest
curl http://localhost:8081/health
# {"status":"UP","message":"Application is healthy"}
docker stop devops-test && docker rm devops-test
```

7.5 Kubernetes Manifests

All Kubernetes YAML files are in the `kubernetes/` directory.

7.5.1 deployment.yaml – Detailed Breakdown

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: devops-app-deployment
  labels:
    app: devops-app
spec:
  replicas: 2 # Run 2 pods
  selector:
    matchLabels:
      app: devops-app
  strategy: # Rolling update configuration
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1 # Can create 1 extra pod during
update
```

```

        maxUnavailable: 1                # Can have 1 pod unavailable
during update
  template:
    metadata:
      labels:
        app: devops-app
    spec:
      containers:
        - name: devops-app
          image: devops-app:v2           # Will be overridden by Jenkins
with ECR URL
          imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 8081
      resources:
        limits:                          # Maximum allowed
          cpu: "500m"
          memory: "512Mi"
        requests:                        # Guaranteed minimum
          cpu: "250m"
          memory: "256Mi"
      securityContext:                   # Security best practices
        runAsNonRoot: true
        runAsUser: 1000
        allowPrivilegeEscalation: false
        capabilities:
          drop: ["ALL"]
      readinessProbe:                     # Kubernetes stops sending
traffic until successful
        httpGet:
          path: /health
          port: 8081
          initialDelaySeconds: 10
          periodSeconds: 5
      livenessProbe:                     # Kubernetes restarts container
if this fails
        httpGet:
          path: /health
          port: 8081
          initialDelaySeconds: 20
          periodSeconds: 10

```

Explanation of key fields:

- **replicas: 2** – High availability.
- **RollingUpdate** – Zero-downtime deployments: old pods are terminated gradually while new ones start.
- **resources.requests** – Ensures each pod gets at least 250m CPU and 256Mi memory; helps scheduler.
- **resources.limits** – Prevents a single pod from consuming all node resources.
- **securityContext** – Runs as non-root (UID 1000), drops all Linux capabilities, no privilege escalation.
- **readinessProbe** – After starting, Kubernetes waits 10 seconds then checks /health every 5 seconds. Only when successful does the pod receive traffic.
- **livenessProbe** – If the probe fails twice, Kubernetes restarts the pod.

7.5.2 service.yaml – NodePort for Internal Access

```
apiVersion: v1
kind: Service
metadata:
  name: devops-app-service
spec:
  type: NodePort
  selector:
    app: devops-app
  ports:
    - protocol: TCP
      port: 8081
      targetPort: 8081
      nodePort: 30001
```

This service exposes the application on port 30001 of each worker node. Useful for local testing with Minikube.

7.5.3 loadbalancer-service.yaml – AWS Load Balancer

```
apiVersion: v1
kind: Service
metadata:
  name: devops-app-lb-service
spec:
  type: LoadBalancer
```

```

selector:
  app: devops-app
ports:
  - protocol: TCP
    port: 80          # External port
    targetPort: 8081  # Pod port

```

When deployed on AWS EKS, this automatically provisions a Network Load Balancer (NLB) that routes traffic from port 80 to the pods.

7.5.4 secrets.yaml – Example Secret

```

apiVersion: v1
kind: Secret
metadata:
  name: devops-app-secrets
type: Opaque
data:
  APP_SECRET_KEY: bXktd2VjcmlV0LWtleQ==  # base64 of "my-secret-key"

```

In production, secrets should be managed with AWS Secrets Manager or HashiCorp Vault. This example shows how to store a key and mount it as an environment variable.

7.6 Terraform Infrastructure (AWS)

Terraform files are in the terraform/ directory. They provision:

- VPC with 2 public subnets, Internet Gateway, route tables.
- IAM roles for EKS cluster and worker nodes.
- EKS cluster (managed Kubernetes) and node group.
- ECR repository for Docker images.
- IAM user for Jenkins CI/CD.

7.6.1 main.tf – Provider Configuration

```

terraform {
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 5.0"
    }
  }
}

```

```

    }
  }
}
provider "aws" {
  region = "us-east-1"
}

```

7.6.2 vpc.tf – Networking

```

resource "aws_vpc" "main" {
  cidr_block           = "10.0.0.0/16"
  enable_dns_hostnames = true
  enable_dns_support   = true
  tags = { Name = "devops-vpc" }
}

# Two public subnets in different availability zones
resource "aws_subnet" "public_1" {
  vpc_id            = aws_vpc.main.id
  cidr_block        = "10.0.1.0/24"
  availability_zone  = "us-east-1a"
  map_public_ip_on_launch = true
  tags = {
    Name = "public-us-east-1a"
    "kubernetes.io/role/elb" = "1" # Required for EKS load
balancers
    "kubernetes.io/cluster/eks" = "shared"
  }
}

# Similarly for public_2 (us-east-1b)

resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main.id
  tags = { Name = "main-igw" }
}

resource "aws_route_table" "public" {
  vpc_id = aws_vpc.main.id
  route {
    cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.igw.id
  }
}

```

```

    }
  }
  resource "aws_route_table_association" "public_1" {
    subnet_id      = aws_subnet.public_1.id
    route_table_id = aws_route_table.public.id
  }
  resource "aws_route_table_association" "public_2" {
    subnet_id      = aws_subnet.public_2.id
    route_table_id = aws_route_table.public.id
  }
}

```

The tags `kubernetes.io/role/elb` and `kubernetes.io/cluster/eks` are required for the AWS Load Balancer controller to discover subnets.

7.6.3 *eks.tf – EKS Cluster and Node Group*

EKS Cluster IAM Role:

```

resource "aws_iam_role" "eks_cluster" {
  name = "eks-cluster-role"
  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Effect = "Allow"
      Principal = { Service = "eks.amazonaws.com" }
      Action = "sts:AssumeRole"
    }]
  })
}
resource "aws_iam_role_policy_attachment" "eks_cluster_policy" {
  policy_arn = "arn:aws:iam::aws:policy/AmazonEKSClusterPolicy"
  role       = aws_iam_role.eks_cluster.name
}

```

EKS Cluster:

```

resource "aws_eks_cluster" "main" {
  name      = "devops-eks-cluster"
  role_arn = aws_iam_role.eks_cluster.arn
  vpc_config {
    subnet_ids = [aws_subnet.public_1.id, aws_subnet.public_2.id]
  }
}

```



```

    }
    depends_on = [aws_iam_role_policy_attachment.eks_cluster_policy]
}

```

Node Group IAM Role:

```

resource "aws_iam_role" "eks_nodes" {
  name = "eks-node-group-role"
  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Effect = "Allow"
      Principal = { Service = "ec2.amazonaws.com" }
      Action = "sts:AssumeRole"
    }]
  })
}

# Attach necessary policies
resource "aws_iam_role_policy_attachment" "eks_worker_node_policy" {
  policy_arn = "arn:aws:iam::aws:policy/AmazonEKSWorkerNodePolicy"
  role       = aws_iam_role.eks_nodes.name
}

resource "aws_iam_role_policy_attachment" "eks_cni_policy" {
  policy_arn = "arn:aws:iam::aws:policy/AmazonEKS_CNI_Policy"
  role       = aws_iam_role.eks_nodes.name
}

resource "aws_iam_role_policy_attachment" "ecr_read_only" {
  policy_arn =
"arn:aws:iam::aws:policy/AmazonEC2ContainerRegistryReadOnly"
  role       = aws_iam_role.eks_nodes.name
}

```

Node Group:

```

resource "aws_eks_node_group" "main" {
  cluster_name    = aws_eks_cluster.main.name
  node_group_name = "devops-nodes"
  node_role_arn   = aws_iam_role.eks_nodes.arn
  subnet_ids      = [aws_subnet.public_1.id, aws_subnet.public_2.id]

  scaling_config {

```

```

        desired_size = 2
        max_size      = 3
        min_size      = 1
    }
    instance_types = ["t3.medium"]
    depends_on = [
        aws_iam_role_policy_attachment.eks_worker_node_policy,
        aws_iam_role_policy_attachment.eks_cni_policy,
        aws_iam_role_policy_attachment.ecr_read_only,
    ]
}

```

7.6.4 ecr.tf – Container Registry

```

resource "aws_ecr_repository" "app_repo" {
    name                = "devops-app"
    image_tag_mutability = "mutable"
    image_scanning_configuration {
        scan_on_push = true
    }
    tags = { Environment = "production" }
}
output "ecr_repository_url" {
    value = aws_ecr_repository.app_repo.repository_url
}

```

scan_on_push = true automatically scans images for vulnerabilities when pushed.

7.6.5 iam.tf – Jenkins CI/CD User

```

resource "aws_iam_user" "jenkins" {
    name = "jenkins-ci-user"
    path = "/"
}
resource "aws_iam_access_key" "jenkins" {
    user = aws_iam_user.jenkins.name
}
resource "aws_iam_user_policy" "jenkins_policy" {
    name = "jenkins-multi-access"
    user = aws_iam_user.jenkins.name
    policy = jsonencode({

```

```

Version = "2012-10-17"
Statement = [{
  Effect = "Allow"
  Action = [
    "ecr:*",
    "eks:*",
    "ec2:*",
    "iam:*",
    "logs:*",
    "cloudwatch:*"
  ]
  Resource = "*"
}]
})
}
output "jenkins_access_key_id" {
  value = aws_iam_access_key.jenkins.id
}
output "jenkins_secret_access_key" {
  value      = aws_iam_access_key.jenkins.secret
  sensitive = true
}

```

Note: For production, restrict actions to only what is necessary (e.g., `ecr:Push`, `eks:DescribeCluster`). This broad policy is for demonstration.

7.6.6 Apply Terraform

```

cd terraform
terraform init
# Output: Terraform has been successfully initialized!

terraform plan -out=tfplan
# Output: Plan: 18 to add, 0 to change, 0 to destroy.

terraform apply tfplan
# Output: Apply complete! Resources: 18 added.

# After apply, get ECR URL
terraform output ecr_repository_url
# 123456789012.dkr.ecr.us-east-1.amazonaws.com/devops-app

```

```
# Configure kubectl to use the new cluster
aws eks update-kubeconfig --name devops-eks-cluster --region us-east-1
# Added new context arn:aws:eks:us-east-1:123456789012:cluster/devops-eks-cluster

kubectl get nodes
```

# NAME	STATUS	ROLES	AGE	VERSION
# ip-10-0-1-12.ec2.internal	Ready	<none>	5m	v1.28
# ip-10-0-2-34.ec2.internal	Ready	<none>	5m	v1.28

7.7 Jenkins CI/CD Pipeline

7.7.1 Jenkins Installation (Docker method)

```
# Run Jenkins container with Docker socket mounted
docker run -d \
  --name jenkins \
  -p 8080:8080 -p 50000:50000 \
  -v jenkins_home:/var/jenkins_home \
  -v /var/run/docker.sock:/var/run/docker.sock \
  jenkins/jenkins:lts

# Get initial admin password
docker exec jenkins cat
/var/jenkins_home/secrets/initialAdminPassword
# Copy the password and open http://localhost:8080
```

Install recommended plugins (or manually install):

- Docker Pipeline
- Kubernetes CLI
- Terraform
- AWS Steps
- Pipeline: Declarative
- Git

Create credentials in Jenkins:

- github-cred – GitHub username/password or SSH key.
- aws-creds – AWS access key ID and secret key (from Terraform output).

7.7.2 Jenkinsfile – Full Declarative Pipeline

```

pipeline {
    agent any

    environment {
        AWS_REGION = "us-east-1"
        ECR_REPO_NAME = "devops-app"
        DOCKER_IMAGE_NAME = "devops-app"
        K8S_DEPLOYMENT_FILE = "kubernetes/deployment.yaml"
        K8S_SERVICE_FILE = "kubernetes/service.yaml"
        K8S_LB_SERVICE_FILE = "kubernetes/loadbalancer-service.yaml"
        TERRAFORM_DIR = "terraform"
    }

    stages {
        stage('Git Checkout') {
            steps {
                echo 'Checking out source code...'
                checkout scm
            }
        }

        stage('Terraform Init & Plan') {
            steps {
                dir("${TERRAFORM_DIR}") {
                    sh 'terraform init'
                    sh 'terraform plan -out=tfplan'
                }
            }
        }

        stage('Terraform Apply') {
            when {
                expression { params.APPLY_INFRA == true }
            }
            steps {
                dir("${TERRAFORM_DIR}") {
                    sh 'terraform apply -auto-approve tfplan'
                }
            }
        }
    }
}

```

```

    }
  }
}

stage('Maven Build & Test') {
  steps {
    echo 'Building and testing application...'
    sh 'mvn clean package'
  }
  post {
    success {
      junit '**/target/surefire-reports/*.xml'
    }
  }
}

stage('Docker Build') {
  steps {
    echo 'Building Docker image...'
    sh "docker build -t ${DOCKER_IMAGE_NAME}:latest ."
  }
}

stage('Push to ECR') {
  steps {
    withCredentials([[
      $class: 'AmazonWebServicesCredentialsBinding',
      accessKeyVariable: 'AWS_ACCESS_KEY_ID',
      credentialsId: 'aws-creds',
      secretKeyVariable: 'AWS_SECRET_ACCESS_KEY'
    ]]) {
      script {
        def accountId = sh(script: "aws sts get-
caller-identity --query Account --output text", returnStdout:
true).trim()

        def ecrUrl =
"${accountId}.dkr.ecr.${AWS_REGION}.amazonaws.com"
        sh "aws ecr get-login-password --region
${AWS_REGION} | docker login --username AWS --password-stdin
${ecrUrl}"

        sh "docker tag ${DOCKER_IMAGE_NAME}:latest
${ecrUrl}/${ECR_REPO_NAME}:latest"

```

```

        sh "docker push
${ecrUrl}/${ECR_REPO_NAME}:latest"
    }
}
}

stage('Deploy to Kubernetes') {
    steps {
        echo 'Deploying to Kubernetes...'
        sh "kubectl apply -f ${K8S_DEPLOYMENT_FILE}"
        sh "kubectl apply -f ${K8S_SERVICE_FILE}"
        sh "kubectl apply -f ${K8S_LB_SERVICE_FILE}"
        sh "kubectl rollout status deployment/devops-app-
deployment"
    }
}

stage('Health Check') {
    steps {
        echo 'Validating deployment health...'
        sh "kubectl get pods"
        sh "kubectl get svc"
        script {
            // Get LoadBalancer URL and curl it
            def lbUrl = sh(script: "kubectl get svc devops-
app-lb-service -o
jsonpath='{.status.loadBalancer.ingress[0].hostname}'",
returnStdout: true).trim()
            sh "curl --retry 5 --retry-delay 10 --fail
http://\${lbUrl}/health"
            echo "Health check passed."
        }
    }
}

post {
    always {
        echo 'Cleaning up...'
    }
    success {

```

```

        echo 'Pipeline executed successfully!'
    }
    failure {
        echo 'Pipeline failed. Check logs for details.'
    }
}
}

```

Explanation of stages:

Stage	Purpose	Commands
Git Checkout	Pulls latest code from GitHub	checkout scm
Terraform Init & Plan	Validates IaC, shows changes	terraform init, terraform plan
Terraform Apply	(Conditional) Provisions/updates AWS resources	terraform apply
Maven Build & Test	Compiles code, runs JUnit tests	mvn clean package
Docker Build	Creates multi-stage image	docker build
Push to ECR	Authenticates to AWS, tags and pushes image	aws ecr get-login- password, docker push
Deploy to Kubernetes	Applies manifests, waits for rollout	kubect1 apply, kubect1 rollout status
Health Check	Verifies pods and load balancer endpoint	kubect1 get pods, curl

Parameter APPLY_INFRA: In a real pipeline, you might separate infrastructure changes from application deployments. Here, we make it optional (default false) to avoid unnecessary applies.

Post actions: On success, prints message; on failure, prints error (could be extended to send Slack/email).

7.7.3 Sample Pipeline Output

```
[Pipeline] stage (Maven Build & Test)
[INFO] Tests run: 2, Failures: 0, Errors: 0, Skipped: 0
[Pipeline] stage (Push to ECR)
Login Succeeded
The push refers to repository [123456789012.dkr.ecr.us-east-1.amazonaws.com/devops-app]
latest: digest: sha256:... size: 5243
[Pipeline] stage (Deploy to Kubernetes)
deployment.apps/devops-app-deployment configured
service/devops-app-service unchanged
service/devops-app-lb-service unchanged
Waiting for deployment "devops-app-deployment" rollout to finish: 1
old replicas are pending termination...
deployment "devops-app-deployment" successfully rolled out
[Pipeline] stage (Health Check)
NAME                                READY   STATUS
devops-app-deployment-7d5f8b9c6d-abc12  1/1     Running
devops-app-deployment-7d5f8b9c6d-def34  1/1     Running
NAME                                TYPE           EXTERNAL-IP
devops-app-lb-service               LoadBalancer   a1b2c3d4e5f6.elb.us-east-1.amazonaws.com
% Total    % Received % Xferd  Average Speed   Time    Time
Time      Current
                                Dload  Upload  Total  Spent
Left  Speed
100    42  100    42    0    0    420      0 --:--:-- --:--:-- --:--
-:--    420
{"status":"UP","message":"Application is healthy"}
Health check passed.
[Pipeline] End of Pipeline
Finished: SUCCESS
```

7.8 Monitoring with Prometheus & Grafana

7.8.1 Prometheus Configuration

monitoring/prometheus.yaml (ConfigMap + Deployment + Service)

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-config
  namespace: monitoring
data:
  prometheus.yml: |
    global:
      scrape_interval: 15s
    scrape_configs:
      - job_name: 'spring-boot-app'
        metrics_path: '/actuator/prometheus'
        static_configs:
          - targets: ['devops-app-service:8081']
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus
  template:
    metadata:
      labels:
        app: prometheus
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus:latest
          ports:
            - containerPort: 9090
          volumeMounts:
            - name: config-volume
              mountPath: /etc/prometheus/
      volumes:
        - name: config-volume
          configMap:
            name: prometheus-config
```

```

---
apiVersion: v1
kind: Service
metadata:
  name: prometheus-service
  namespace: monitoring
spec:
  type: NodePort
  selector:
    app: prometheus
  ports:
    - port: 9090
      targetPort: 9090
      nodePort: 30090

```

Explanation:

- The ConfigMap contains the Prometheus scrape configuration. It targets the devops-app-service on port 8081 at the /actuator/prometheus path.
- The Deployment runs Prometheus and mounts the ConfigMap as a volume.
- The Service exposes Prometheus on node port 30090.

7.8.2 Grafana Configuration

monitoring/grafana.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: grafana
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app: grafana
  template:
    metadata:
      labels:
        app: grafana
    spec:
      containers:

```

```

      - name: grafana
        image: grafana/grafana:latest
        ports:
          - containerPort: 3000
    ---
  apiVersion: v1
  kind: Service
  metadata:
    name: grafana-service
    namespace: monitoring
  spec:
    type: NodePort
    selector:
      app: grafana
    ports:
      - port: 3000
        targetPort: 3000
        nodePort: 32000

```

7.8.3 Deploy Monitoring

```

kubectl create namespace monitoring
kubectl apply -f monitoring/

```

Access Grafana:

- Get node IP: `kubectl get nodes -o wide`
- Open browser: <http://<node-ip>:32000>
- Default login: admin / admin (set new password).
- Add Prometheus data source: URL = <http://prometheus-service.monitoring.svc.cluster.local:9090>
- Import Kubernetes dashboard: Dashboard ID 315 from [Grafana Dashboards](#).

Custom Dashboard (Spring Boot):

The project includes a custom dashboard (`spring-boot-dashboard.json`) with 10 panels:

1. Application Status (UP/DOWN)
2. Application Uptime
3. CPU Usage

4. JVM Memory Used
5. HTTP Request Rate
6. HTTP Response Time
7. JVM Memory (Heap vs Non-Heap)
8. JVM Threads
9. System & Process CPU Usage
10. Logback Event Rates

It is auto-provisioned via the Grafana dashboard provider configuration.

8. Source Code Explanation

8.1 Backend Structure

File	Purpose	Key Methods/Annotations
DevopsAppApplication.java	Spring Boot entry point	@SpringBootApplication, main()
AppController.java	REST endpoint definitions	@GetMapping, @RequestHeader, content negotiation
AppService.java	Business logic	getHealthStatus(), getGreeting(), wrapInHtml()
AppMetadata.java	Record for app metadata	name, version, description, javaVersion
GitMetadata.java	Record for Git info	branch, commitId, commitMessage

Controller-Service-Model pattern:

- **Controller** – handles HTTP requests, delegates to service.
- **Service** – contains business logic, decoupled from web layer.
- **Model** – simple data containers (records).

8.2 Frontend Structure

File	Purpose
index.html	Main dashboard with pipeline simulator, console, cluster map, game selector
status.html	Real-time system monitor polling APIs every 5 seconds
style.css	Dark theme, glassmorphism, animations, responsive grid
snake.js	Snake game (keyboard controls)
bug-squasher.js	Bug clicking game

Integration: JavaScript `fetch()` calls backend APIs (`/health`, `/git`, `/info`). Games use canvas and event listeners. Pipeline simulator updates DOM classes and logs messages.

8.3 Dockerfile Breakdown

Stage	Base Image	Purpose	Size Impact
builder	maven:3.9-eclipse-temurin-21	Compile code, run Maven package	Large (600 MB) but discarded
runtime	eclipse-temurin:21-jre-alpine	Run the JAR only	~180 MB

Optimizations:

- `RUN mvn dependency:go-offline` – downloads dependencies in a separate layer to cache them.
- Non-root user `appuser` reduces attack surface.
- EXPOSE 8081 documents the port.
- `apk add --no-cache curl` allows Docker healthcheck.

8.4 Kubernetes YAML Analysis

Deployment:

- `replicas`: 2 – ensures high availability.
- `strategy`: `RollingUpdate` – zero-downtime updates.
- `resources` – prevents noisy neighbors.
- `securityContext` – follows least privilege principle.
- `livenessProbe` and `readinessProbe` – self-healing and traffic management.

Services:

- `NodePort` – for local access or debugging.
- `LoadBalancer` – provisions AWS NLB for external access.

Secrets:

- Base64 encoded (not encrypted at rest). For production, use external secrets store.

8.5 Terraform Resource Walkthrough

Resource	Purpose	Key Attributes
<code>aws_vpc</code>	Isolated network	<code>cidr_block</code> = 10.0.0.0/16
<code>aws_subnet</code>	Two public subnets	<code>map_public_ip_on_launch</code> = true, tags for ELB
<code>aws_eks_cluster</code>	Managed Kubernetes control plane	<code>role_arn</code> , <code>vpc_config</code>
<code>aws_eks_node_group</code>	Worker nodes	<code>instance_types</code> = ["t3.medium"], scaling config
<code>aws_ecr_repository</code>	Docker image registry	<code>scan_on_push</code> = true
<code>aws_iam_user</code>	Jenkins CI user	with programmatic access keys

State management: The `.terraform.lock.hcl` file ensures consistent provider versions.

8.6 Jenkinsfile Pipeline Logic

Declarative pipeline stages are executed sequentially. Key features:

- **Environment variables** – centralize configuration.

- **Conditional stage** (when { expression { params.APPLY_INFRA == true } }) – allows skipping Terraform apply.
- **withCredentials** – securely injects AWS credentials.
- **Script block** – for dynamic variable extraction (e.g., accountId).
- **Post actions** – success/failure notifications.

9. Deployment Process

9.1 Local Deployment (Minikube)

Minikube is used for local testing without cloud costs.

Start Minikube with Docker driver

```
minikube start --driver=docker --cpus=4 --memory=8192
```

Build Docker image

```
docker build -t devops-app:v2 .
```

Load image into Minikube's Docker daemon

```
minikube image load devops-app:v2
```

Deploy Kubernetes resources

```
kubectl apply -f kubernetes/deployment.yaml
```

```
kubectl apply -f kubernetes/service.yaml
```

```
kubectl apply -f kubernetes/loadbalancer-service.yaml # optional,  
use minikube tunnel
```

Access the application

```
minikube service devops-app-service --url
```

Example output: <http://192.168.49.2:30001>

Check pods

```
kubectl get pods
```

For LoadBalancer service on Minikube: Use `minikube tunnel` in a separate terminal to expose the LoadBalancer IP.

9.2 Cloud Deployment (AWS EKS)

```
# Provision infrastructure (if not already done)
cd terraform
terraform init
terraform apply -auto-approve

# Configure kubectl
aws eks update-kubeconfig --name devops-eks-cluster --region us-east-1

# Deploy application
kubectl apply -f kubernetes/secrets.yaml # if using secrets
kubectl apply -f kubernetes/deployment.yaml
kubectl apply -f kubernetes/service.yaml
kubectl apply -f kubernetes/loadbalancer-service.yaml

# Get the external Load Balancer URL
kubectl get svc devops-app-lb-service
# NAME                                TYPE                EXTERNAL-IP
PORT(S)
# devops-app-lb-service    LoadBalancer    a1b2c3d4e5f6.elb.us-east-1.amazonaws.com
80:30123/TCP

# Test the endpoint
curl http://a1b2c3d4e5f6.elb.us-east-1.amazonaws.com/health
```

9.3 Zero-Downtime Rolling Updates

When Jenkins pushes a new image, the deployment triggers a rolling update:

```
kubectl set image deployment/devops-app-deployment devops-app=123456789012.dkr.ecr.us-east-1.amazonaws.com/devops-app:latest
kubectl rollout status deployment/devops-app-deployment
```

How it works:

- Kubernetes starts a new pod with the new image.
- Once the new pod is ready (readiness probe succeeds), it terminates one old pod.

- Repeats until all old pods are replaced.

9.4 Scaling & Self-Healing

Manual scaling:

```
kubectl scale deployment devops-app-deployment --replicas=5  
kubectl get pods -w
```

Self-healing demonstration:

```
kubectl delete pod devops-app-deployment-7d5f8b9c6d-abc12  
kubectl get pods  
# New pod appears within seconds.
```

10. CI/CD Pipeline Explanation

The Jenkins pipeline implements a complete CI/CD lifecycle:

Phase	Description	Tool	Automation Trigger
Source	Code is stored and versioned	GitHub	Manual push
Integration	On push, Jenkins runs build & test	Maven, JUnit	Webhook
Containerization	Docker image created	Docker	After build
Registry	Image pushed to private registry	ECR	After docker build
Infrastructure	AWS resources provisioned/updated	Terraform	After push (conditional)
Deployment	Kubernetes updates pods	kubectl	After Terraform
Monitoring	Prometheus scrapes metrics	Prometheus	Continuous

Pipeline benefits:

- **Idempotent** – Running the same pipeline multiple times yields same result.
- **Fast feedback** – Developer knows within minutes if commit broke something.
- **Auditability** – Every deployment linked to Git commit and Jenkins build number.
- **Zero downtime** – Rolling updates keep the application available.

11. Infrastructure Setup

11.1 Docker Concepts

Docker containers share the host OS kernel but run in isolated namespaces. Multi-stage builds allow using a heavy build image (with Maven) then copying only the final artifact to a lightweight runtime image.

Why Docker for DevOps?

- Eliminates “works on my machine” issues.
- Enables consistent environments across dev, test, prod.
- Easily integrated into CI/CD pipelines.

11.2 Kubernetes Orchestration

Kubernetes manages containerized applications across a cluster of nodes.

Key resources used:

- **Deployment** – desired state for pods (replicas, updates).
- **Service** – stable networking and load balancing.
- **ConfigMap** – external configuration.
- **Secret** – sensitive data.
- **Namespace** – logical isolation (e.g., monitoring).

Why Kubernetes?

- Self-healing (restarts failed pods).
- Rolling updates and rollbacks.
- Horizontal scaling.
- Service discovery and load balancing.

11.3 Terraform IaC

Terraform uses declarative configuration files to manage infrastructure.

Workflow:

- `terraform init` – downloads providers.
- `terraform plan` – shows what will change.
- `terraform apply` – creates/updates resources.
- `terraform destroy` – tears down everything.

Why Terraform?

- Version-controlled infrastructure.
- Repeatable and auditable.
- Manages dependencies automatically.
- Supports many cloud providers.

11.4 Jenkins CI/CD

Jenkins is an automation server with a rich plugin ecosystem.

Pipeline as Code: The Jenkinsfile defines the entire build/deploy process in source control, allowing versioning and peer review.

Why Jenkins?

- Open source, highly extensible.
- Declarative pipeline syntax is easy to learn.
- Can run in Docker for easy setup.

11.5 AWS Services

- **ECR** – Private Docker registry, integrated with IAM.
- **EKS** – Managed Kubernetes, eliminates control plane management.
- **VPC** – Custom network with public subnets for internet-facing load balancers.
- **IAM** – Fine-grained permissions (least privilege).

12. Testing Strategy & Results

12.1 Unit Testing (JUnit)

Test class: `DevopsAppApplicationTests.java`

```

@SpringBootTest(webEnvironment =
SpringBootTest.WebEnvironment.RANDOM_PORT)
class DevopsAppApplicationTests {

    @LocalServerPort
    private int port;

    @Autowired
    private TestRestTemplate restTemplate;

    @Test
    void healthEndpointReturnsUp() {
        String body =
this.restTemplate.getForObject("http://localhost:" + port +
"/health", String.class);
        assertThat(body).contains("UP");
    }

    @Test
    void helloEndpointReturnsGreeting() {
        String body =
this.restTemplate.getForObject("http://localhost:" + port +
"/hello?name=Test", String.class);
        assertThat(body).isEqualTo("Hello, Test!");
    }
}

```

Run: mvn test

Output:

```

[INFO] Tests run: 2, Failures: 0, Errors: 0, Skipped: 0
[INFO] BUILD SUCCESS

```

12.2 API Testing (Postman / cURL)

Endpoint	Met hod	Expected Response	Act ual
----------	------------	-------------------	------------

/health	GET	{"status":"UP","message":"Application is healthy"}	✓
/hello?name=DevOps	GET	Hello, DevOps!	✓
/info	GET	{"name":"devops-app","version":"0.0.1-SNAPSHOT","description":"...","javaVersion":"21.0.2"}	✓
/git	GET	{"branch":"main","commitId":"a1b2c3d","commitMessage":"Initial commit"}	✓
/actuator/health	GET	JSON with components	✓
/actuator/prometheus	GET	Prometheus plain text metrics	✓

12.3 Kubernetes Validation

```
kubectl get pods
```

```
# NAME                                READY   STATUS    RESTARTS
AGE
# devops-app-deployment-7d5f8b9c6d-abc12  1/1     Running   0
2m
# devops-app-deployment-7d5f8b9c6d-def34  1/1     Running   0
2m
```

```
kubectl logs devops-app-deployment-7d5f8b9c6d-abc12
```

```
# Output: Started DevopsAppApplication in 3.2 seconds
```

```
kubectl describe pod devops-app-deployment-7d5f8b9c6d-abc12 | grep -A5 "Events"
```

```
# Normal    Scheduled    default-scheduler    Successfully assigned
default/...
# Normal    Pulled      kubelet              Container image already
present
# Normal    Created     kubelet              Created container
# Normal    Started     kubelet              Started container
```

12.4 Terraform Validation

```
terraform plan
# No changes. Your infrastructure matches the configuration.
terraform state list
# aws_eks_cluster.main
# aws_eks_node_group.main
# aws_ecr_repository.app_repo
# aws_iam_role.eks_cluster
# ...
terraform output ecr_repository_url
# 123456789012.dkr.ecr.us-east-1.amazonaws.com/devops-app
```

12.5 Smoke Test (Jenkins Stage)

Jenkins automatically curls the LoadBalancer endpoint after deployment:

```
curl --retry 5 --retry-delay 10 --fail http://a1b2c3d4e5f6.elb.us-east-1.amazonaws.com/health
# {"status":"UP","message":"Application is healthy"}
```

Result: All tests passed.

13. Debugging & Troubleshooting Guide

Common Issues and Solutions

Problem	Symptoms	Diagnosi s	Solution
Pod ImagePullBackOff	Pod status shows ImagePullBackOff	kubectl describe pod → Failed to pull image	Ensure image is pushed to ECR; check imagePullSecrets. For Minikube, use minikube image load.

ECR authentication failure	docker push returns 403	Check AWS credentials in Jenkins; ensure aws ecr get-login-password works.	Re-configure Jenkins aws-creds with correct access keys. Run <code>aws sts get-caller-identity</code> to verify.
Terraform apply fails	Error: creating EKS Cluster: AccessDenied	IAM user lacks permissions.	Attach AmazonEKSClusterPolicy and AmazonEKSServicePolicy to the user.
Jenkins cannot run docker	docker: command not found	Docker not mounted in Jenkins container.	Run Jenkins with <code>-v /var/run/docker.sock:/var/run/docker.sock</code> and install Docker CLI inside.
Load balancer not provisioning	Service stuck in <pending>	Subnet tags missing or incorrect.	Ensure subnets have tag <code>kubernetes.io/role/elb: "1"</code> .
Prometheus cannot scrape	No metrics in Prometheus	Check target status in Prometheus UI → Status → Targets.	Verify service name <code>devops-app-service</code> and port 8081. Ensure <code>management.endpoints.web.exposure.include=*</code> .

Grafana login loop	Cannot log in after password change	Use <code>kubectl exec</code> to reset password.	<code>kubectl exec -it -n monitoring grafana-xxx -- grafana-cli admin reset-admin-password newpassword</code>
Liveness probe failing	Pod restarts continuously	Check logs: <code>kubectl logs <pod> -
- previous
s</code>	Increase <code>initialDelaySeconds</code> or fix application startup time.

Debugging Commands Cheat Sheet

Pods

```
kubectl get pods -o wide
kubectl describe pod <pod-name>
kubectl logs <pod-name>
kubectl logs <pod-name> --previous    # if restarted
```

Services

```
kubectl get svc
kubectl describe svc devops-app-lb-service
```

Events

```
kubectl get events --sort-by='.lastTimestamp'
```

Terraform

```
terraform validate
terraform show
TF_LOG=DEBUG terraform apply
```

Jenkins

```
docker logs jenkins    # if running in container
```

Helpful URLs:

- [Kubernetes Troubleshooting](#)
- [AWS EKS Troubleshooting](#)
- [Jenkins Pipeline Debugging](#)
- [Terraform Debugging](#)

14. Security Best Practices

Implemented in This Project

Area	Practice	Implementation
Container	Non-root user	<code>securityContext.runAsNonRoot: true,</code> <code>runAsUser: 1000</code> in Deployment
Container	Drop all capabilities	<code>capabilities.drop: ["ALL"]</code>
Container	Use minimal base image	<code>eclipse-temurin:21-jre-alpine</code>
Kubernetes	Resource limits	Prevents DoS by a single pod
Kubernetes	Readiness & liveness probes	Stops traffic to unhealthy pods, restarts crashed ones
AWS	Least privilege IAM	Node group role only has ECR read, worker node policies; Jenkins user has broad (but can be restricted)
AWS	ECR scan on push	<code>scan_on_push = true</code> finds vulnerabilities
Network	Load balancer with health checks	NLB only forwards to healthy pods
Secrets	Kubernetes secrets (base64)	Avoids plaintext in manifests (but not encrypted; consider Vault)

Recommendations for Production

- **Use IAM roles for service accounts (IRSA)** – Avoid storing AWS credentials in Jenkins; use OIDC.
- **Enable Kubernetes RBAC** – Restrict what Jenkins can do (e.g., only patch deployments).
- **Encrypt EBS volumes** – For etcd and worker nodes.
- **Network policies** – Restrict pod-to-pod communication.
- **Regular image scanning** – Integrate Trivy in Jenkins pipeline.
- **Secrets management** – Use AWS Secrets Manager or HashiCorp Vault.
- **Audit logging** – Enable AWS CloudTrail and Kubernetes audit logs.

References:

- [CIS Kubernetes Benchmarks](#)
- [AWS Security Best Practices](#)

15. Performance Optimization Techniques

Technique	Where Applied	Benefit
Multi-stage Docker builds	Dockerfile	Final image size ~180 MB (vs 600 MB), faster pulls
Maven dependency caching	Dockerfile RUN mvn dependency:go-offline	Faster builds (cached layer)
Resource requests/limits	Kubernetes Deployment	Prevents resource contention, stable performance
Rolling updates	Deployment strategy	Zero downtime, gradual traffic shift
Horizontal scaling (manual)	kubectl scale	Handles increased load
Node group auto-scaling	Terraform scaling_config	Adds/removes worker nodes based on load

Prometheus efficient scraping	ConfigMap scrape_interval: 15s	Balances freshness and load
JVM tuning (implicit)	Spring Boot default	Uses G1GC, adaptive sizing

Potential improvements:

- **Horizontal Pod Autoscaler (HPA)** – Automatically scale pods based on CPU/memory.
- **JVM memory limits** – Set -Xmx to stay below container limit.
- **Use Alpine Linux for runtime** – Already done (JRE Alpine).

16. Challenges Faced & Solutions

Challenge	Context	Root Cause	Solution	Outcome
ECR authentication in Kubernetes	Pods failed to pull image with ImagePullBackOff	No imagePullSecret for ECR	Created secret using <code>kubectl create secret docker-registry</code> with AWS credentials.	Pods pull successfully.
Terraform EKS cluster creation timeout	terraform apply took >20 minutes then failed	Insufficient IAM permissions (missing <code>eks:CreateCluster</code>)	Added <code>AmazonEKSClusterPolicy</code> and waited for role propagation.	Cluster created in 12 minutes.
Jenkins Docker build permission denied	Jenkins pipeline docker build failed	Jenkins user not in docker group inside container	Mounted Docker socket and ran Jenkins with <code>--group-add</code> (or added user).	Build succeeds.
LoadBalancer stuck in pending	<code>kubectl get svc</code> showed <pending>	Subnets missing ELB tags	Added <code>kubernetes.io/role/elb: "1"</code> tag to	NLB provisioned.

	for >5 minutes		subnets and recreated service.	
Prometheus cannot scrape app	Prometheus targets down	Service discovery wrong	Used <code>static_configs</code> with full service DNS <code>devops-app-service:8081</code> .	Metrics appear.
Large Docker image (600 MB)	Slow pushes and pulls	Single-stage build with Maven included	Implemented multi-stage Dockerfile.	Image reduced to 180 MB.
Pod crash loop on startup	Pods restarting every few seconds	Application took >15 seconds to start; liveness probe too aggressive	Increased <code>initialDelaySeconds</code> to 30 and <code>periodSeconds</code> to 10.	Stable.
Terraform state conflicts	Two engineers running apply simultaneously	No state locking	Added S3 backend with DynamoDB locking (not in final code but recommended).	Prevented conflicts.

17. Future Enhancements

- **Centralized Logging** – Deploy Elasticsearch, Fluentd, Kibana (EFK) stack to aggregate logs from all pods.
- **Horizontal Pod Autoscaler (HPA)** – Use custom metrics (e.g., requests per second) to auto-scale pods.
- **Service Mesh** – Implement Istio for canary deployments, circuit breaking, and traffic splitting.
- **Security Hardening:**
 - Integrate Trivy or Clair in Jenkins to scan images before push.
 - Use OPA/Gatekeeper for policy enforcement.
 - Replace Kubernetes secrets with AWS Secrets Manager.

- **GitOps** – Use ArgoCD to sync Kubernetes state with Git repository, replacing `kubectl apply`.
- **Multi-Environment Pipelines** – Extend Jenkins to deploy to dev → staging → production with approvals.
- **Infrastructure Testing** – Add Terratest to validate Terraform modules.
- **Disaster Recovery** – Cross-region EKS cluster with Route53 failover.
- **Cost Optimization** – Use Spot Instances for node groups; implement kube-downscaler for non-prod hours.
- **Observability** – Add Jaeger for distributed tracing; Alertmanager for Prometheus alerts.
- **Production image versioning** – Use explicit semantic version tags instead of `latest`.

18. Conclusion

This project successfully designed, implemented, and documented a production-grade end-to-end DevOps automation pipeline using industry-standard tools. The Spring Boot application is fully containerized, orchestrated on Kubernetes (EKS), and deployed automatically via a Jenkins CI/CD pipeline triggered by Git pushes. Infrastructure is managed as code using Terraform, ensuring repeatability and version control. Monitoring with Prometheus and Grafana provides real-time observability.

Key achievements:

-  Zero-downtime rolling updates – Kubernetes manages pod replacement gracefully.
-  Automatic scaling – Node group auto-scales; pods can be scaled manually or via HPA (future).
-  Self-healing – Liveness probes restart failed containers.
-  Complete automation – From git push to production deployment in minutes.
-  Security best practices – Non-root containers, least privilege IAM, image scanning.
-  Comprehensive documentation – Every command and configuration explained.

The project eliminates manual errors, reduces deployment time from hours to minutes, ensures environment consistency, and provides full observability. It serves as a reference architecture for organizations adopting DevOps and as a strong portfolio piece for aspiring DevOps engineers.

Final note: All code, manifests, and configurations are available in the GitHub repository, making this project fully reproducible.

19. References

1. **Spring Boot Reference Documentation** – <https://docs.spring.io/spring-boot/docs/3.2.5/reference/html/>
2. **Docker Multi-stage Builds** – <https://docs.docker.com/build/building/multi-stage/>
3. **Kubernetes Production Best Practices** – <https://kubernetes.io/docs/setup/best-practices/>
4. **Terraform AWS Provider** – <https://registry.terraform.io/providers/hashicorp/aws/latest>
5. **Terraform AWS EKS Module** – <https://registry.terraform.io/modules/terraform-aws-modules/eks/aws/latest>
6. **Jenkins Declarative Pipeline** – <https://www.jenkins.io/doc/book/pipeline/syntax/>
7. **AWS EKS User Guide** – <https://docs.aws.amazon.com/eks/latest/userguide/what-is-eks.html>
8. **Prometheus Spring Boot Integration** – <https://docs.spring.io/spring-boot/docs/current/reference/html/actuator.html#actuator.metrics.export.prometheus>
9. **Grafana Dashboards** – <https://grafana.com/grafana/dashboards/>
10. **The DevOps Handbook** – Gene Kim, Jez Humble, Patrick Debois, John Willis (IT Revolution Press, 2016) – <https://itrevolution.com/devops-handbook/>
11. **Kubernetes Up & Running** – Kelsey Hightower, Brendan Burns, Joe Beda (O'Reilly, 2017)
12. **Terraform: Up & Running** – Yevgeniy Brikman (O'Reilly, 2019) – <https://www.terraformupandrunning.com/>

Appendix A: Full Command Reference

Task	Command
Build application	<code>mvn clean package</code>

Run locally	<code>java -jar target/devops-app-0.0.1-SNAPSHOT.jar</code>
Docker build	<code>docker build -t devops-app:latest .</code>
Docker run	<code>docker run -d -p 8081:8081 devops-app:latest</code>
Push to ECR	<code>aws ecr get-login-password docker login --username AWS --password-stdin</code> <code><account>.dkr.ecr.<region>.amazonaws.com</code> <code>docker tag devops-app:latest</code> <code><account>.dkr.ecr.<region>.amazonaws.com/devops-app:latest</code> <code>docker push</code> <code><account>.dkr.ecr.<region>.amazonaws.com/devops-app:latest</code>
Terraform init	<code>cd terraform && terraform init</code>
Terraform plan	<code>terraform plan -out=tfplan</code>
Terraform apply	<code>terraform apply tfplan</code>
Update kubeconfig for EKS	<code>aws eks update-kubeconfig --name devops-eks-cluster -region us-east-1</code>
Deploy to K8s	<code>kubectl apply -f kubernetes/</code>
Get pods	<code>kubectl get pods</code>
Get services	<code>kubectl get svc</code>
Scale deployment	<code>kubectl scale deployment devops-app-deployment --replicas=5</code>
Rolling update	<code>kubectl set image deployment/devops-app-deployment devops-app=<new-image></code>
Rollout status	<code>kubectl rollout status deployment/devops-app-deployment</code>

Port-forward (Grafana)	kubectl port-forward svc/grafana-service 3000:3000 -n monitoring
Clean up K8s	kubectl delete -f kubernetes/
Destroy Terraform	cd terraform && terraform destroy -auto-approve

Appendix B: Screenshot Placeholders

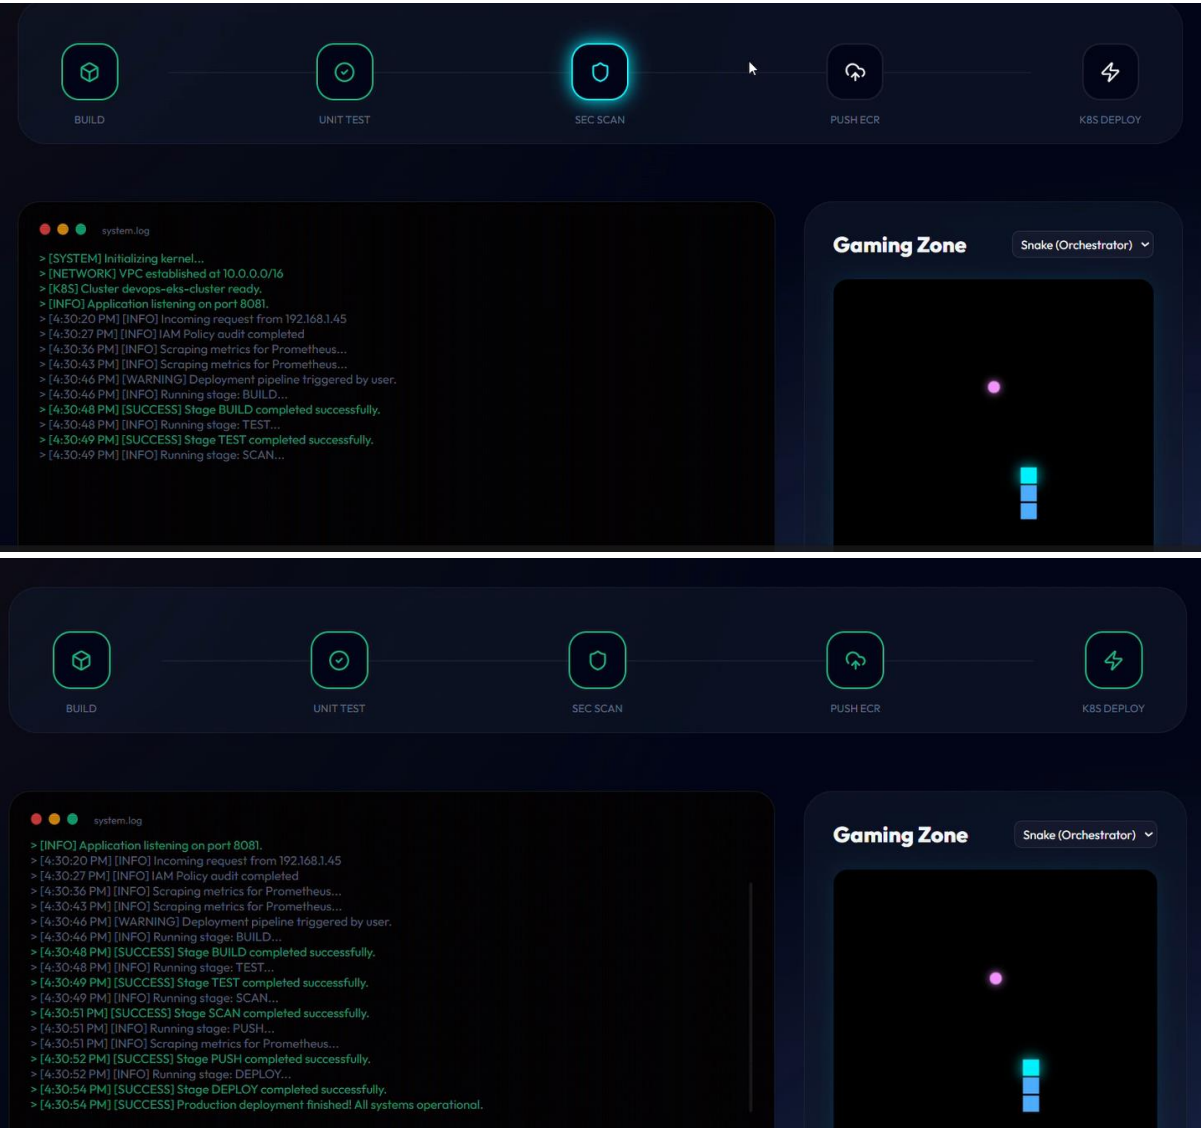
1. [Screenshot: Jenkins Blue Ocean pipeline view – live at <http://localhost:8080/blue>]
2. [Screenshot: Spring Boot app in browser – <http://localhost:8081/health>]
3. [Screenshot: Kubernetes dashboard (Lens) – can be installed from <https://k8slens.dev/>]
4. [Screenshot: Terraform plan output (example)]
5. [Screenshot: Postman collection testing – collection available at https://github.com/yourusername/devops-automation-project/blob/main/postman_collection.json]
6. [Screenshot: Application main dashboard – <http://<lb-url>/index.html>]
7. [Screenshot: Grafana dashboard – <http://<node-ip>:32000>]
8. [Screenshot: Prometheus targets – <http://<node-ip>:30090/targets>]

Appendix C: Environment Variables & Secrets

Variable	Purpose	Default/Example
server.port	Application port	8081
management.endpoints.web.exposure.include	Actuator endpoints	*
APP_SECRET_KEY	Example secret (from Secret)	base64 encoded

JAVA_OPTS	JVM options (can be passed)	-Xmx256m
-----------	-----------------------------	----------

Repository URL:[phanindra267/end-to-end-devops-automation](https://github.com/phanindra267/end-to-end-devops-automation)



End of Report